




Binary Search Trees (BST)

Reminder:
Basics of
Binary Search Trees

Index




Record

record

Reg. #	Name	Home address	Phone #	Student id	Awards	...
Education history				

하나의 record는 { 한 block의 일부,
한 block 전체, 또는
여러 block } 을 점유

Blocks



<그림. from Wikipedia>

Sector: 512, 4K, ... bytes

Disk block: multiple of sectors

File block: multiple of disk blocks

- * OS에서 file system block size 지정 가능
- * 한 컴퓨터가 여러 file system을 가질 수 있다


Record, Key, Search Tree

- A record
 - Contains all information for an object
 - e.g., human record
 - <resident registration #주민번호, name, home address, phone #s, education history, income, family members, ... > ← a collection of such fields
- Field
 - each unit information in a record
 - e.g., in the above record, resident registration #, name, home address, ...
- Search key or Key
 - A field which can uniquely identifies each record
 - A key may consist of a field or more
- Search tree
 - A key in each node
 - An index for keys and corresponding record locations


Binary Search Trees

- A key in each node. All keys are distinct.
- Root node in the top level
- Each node has at most two children.
- The key of a node is greater than all the keys of its left subtree, and smaller than all the keys of its right subtree.

Examples of Binary Search Trees





(a)




(b)

Examples of Subtrees



(b) Node r 's left subtree



(c) Node r 's right subtree

(Static) Optimal Binary Search Tree

Possible to construct optimal binary search tree.

See <Dynamic Programming>.

Reminder: Search in Binary Search Trees

search(t, x):

◀ t : root node, x : key


if ($t = \text{NIL}$ **or** $t.\text{key} = x$) **then return** t

if ($x < t.\text{key}$)


then return **search**($t.\text{left}, x$)

else return **search**($t.\text{right}, x$)

In real fields, a node contains
the **key** field and the
corresponding **record location**



Recursive View of Search



Reminder: Insertion in Binary Search Trees

insert(t, x):

◀ t : root node, x : key to insert

if ($t = \text{NIL}$)

$r.\text{key} \leftarrow x$

◀ r : new node

return r


if ($x < t.\text{key}$)

$t.\text{left} \leftarrow \text{insert}(t.\text{left}, x)$

return t

else $t.\text{right} \leftarrow \text{insert}(t.\text{right}, x)$


return t




Examples of Insertion




(a)




(b)




(c)



(d)



(e)



(f)

Reminder: Deletion in Binary Search Trees

t : root node

r : node to deleted

There are three cases

- Case 1 : r is a leaf node
- Case 2 : r has only one child
- Case 3 : r has two children

Reminder: Deletion in Binary Search Trees

sketch_delete(t, r):

◀ t : root node, r : node to delete

if (r is a leaf node)

◀ Case 1

Just throw away r

else if (r has only one child)

◀ Case 2


Let r 's parent links to the (only) child of r

else


◀ Case 3

Remove the minimum node s of r 's right subtree,
and copy the key of s to node r

Example: Case 1




(a) *r* has no child




(b) Simply throw away *r*


Example: Case 2



(a) r has only one child




(b) Remove r




(c) Put r 's (only) child at r 's location


Example: Case 3




(a) Find r 's inorder successor s




(b) Remove r (imaginary removal)





(c) Move s to r's location
(copy s to r)




(d) Move s's (only) child to s's location

Efficiency of Binary Search Trees

IPL: Internal Path Length

Sum of depths of all the nodes

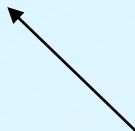


Theorem

The IPL of a binary search tree made at random is $O(n \log n)$ on average (Assume every permutation of the input sequence is equally likely)

<Proof>

Next page




Equivalently, a sequence of n inserts into an empty binary search tree takes $O(n \log n)$ on average (Assume every permutation of the input sequence is equally likely)

✓ **Meaning:** Average search time for a key is $O(\log n)$.

<Proof>

$D(n)$: the average IPL (Internal Path Length) of a binary tree with n nodes.

Clearly $D(0) = 0$, $D(1) = 1$.



$$\begin{aligned} D(n) &= \frac{1}{n} \sum_{k=1}^n [D(k-1) + (k-1) + D(n-k) + (n-k) + 1] \\ &= \frac{2}{n} \sum_{k=0}^{n-1} D(k) + n \end{aligned}$$

Assume that $\exists c > 0$ s.t. $D(k) \leq ck \log k \ \forall k < n$.

Then, we verify that $D(n) \leq cn \log n$ (i.e., $D(n) = O(n \log n)$)

$$\begin{aligned}
D(n) &= \frac{2}{n} \sum_{k=0}^{n-1} D(k) + n \\
&= \frac{2}{n} \sum_{k=2}^{n-1} D(k) + \Theta(n) \quad \swarrow D(0), D(1) \text{ absorbed} \\
&\leq \frac{2}{n} \sum_{k=2}^{n-1} ck \log k + \Theta(n) \\
&\leq \frac{2}{n} \int_1^n cx \log x \, dx + \Theta(n) \\
&= \frac{2c}{n} \left(\left[\frac{1}{2} x^2 \log x \right]_1^n - \left[\frac{1}{4} x^2 \right]_1^n \right) + \Theta(n)
\end{aligned}$$

$$= \frac{2c}{n} \left(\frac{1}{2} n^2 \log n - \frac{1}{4} n^2 + \frac{1}{4} \right) + \Theta(n)$$

$$= cn \log n - \frac{cn}{2} + \underbrace{\frac{c}{2n}}_{\text{absorbed}} + \Theta(n)$$

$$= cn \log n - \frac{cn}{2} + \Theta(n)$$

$$\leq cn \log n$$

We can choose $c > 0$ s.t. $\frac{cn}{2}$ dominates $\Theta(n)$

$$\therefore D(n) = O(n \log n)$$

Balanced Binary Search Trees


Reminder: AVL Tree

Devised by Adelson-Velskii and Landis

A balanced search tree

such that

the heights(depths) of the left and right subtrees of any node
differ by **at most 1**




AVL Tree

Covered in <Data Structures>


Four Types of Repairs

1. Type LL: Right rotation




Four Types of Repairs

2. Type LR: Left rotation then right rotation (conversion to type LL)




Four Types of Repairs

Another Instance of Type LR



Four Types of Repairs


3. Type RR: Left rotation




Four Types of Repairs

4. Type RL: Right rotation then left rotation (conversion to type RR)


* LL과 RR, LR과 RL은
각각 symmetric



Type RL



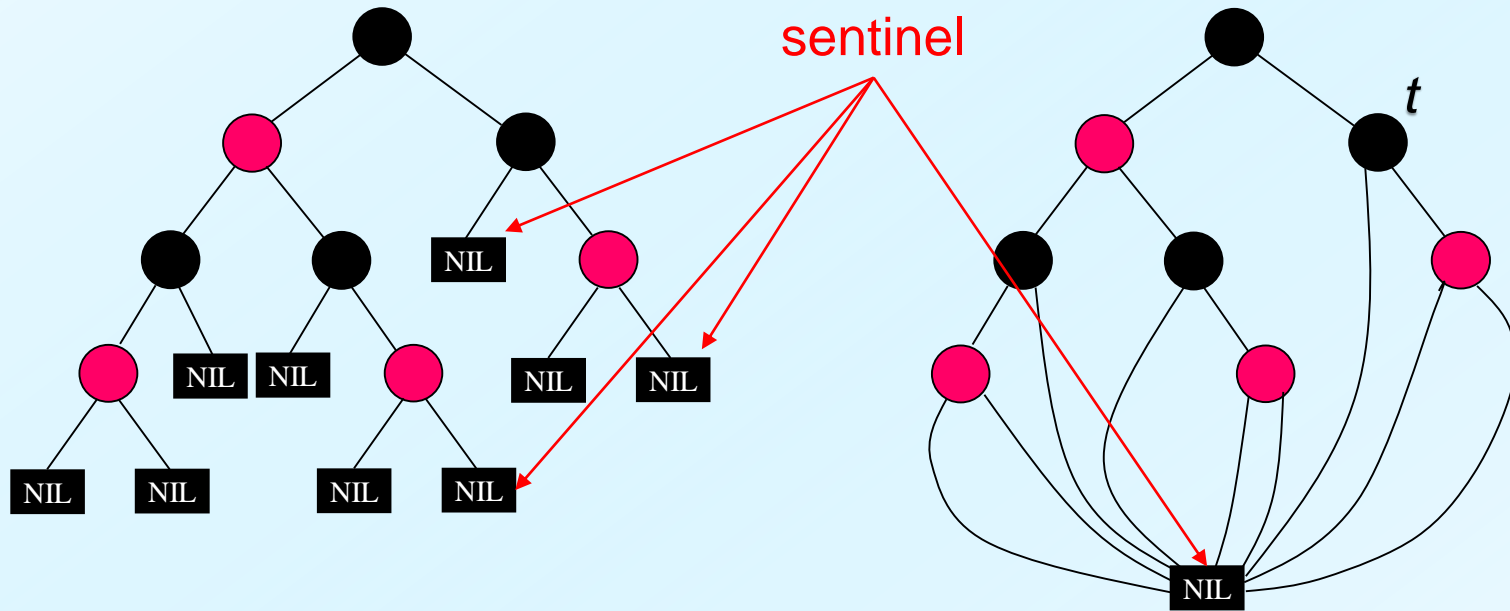
Type RR



Red-Black Trees

- Every node in the search tree has a color: red or black.
 - It has to satisfy the following properties
(**red-black** properties = **RB** properties):
 - ① Every leaf is **black**
 - ② If a node is **red**, its children should be **black** (no two consecutive **reds**)
 - ③ In any path from the root to a leaf, the # of **black** nodes on the path is the same (**black height**)
- ✓ Here, a leaf is not a general leaf node.
Every **null** reference links to the **NIL** leaf node(sentinel).
- ✓ 보통은 root가 **black**이라는 성질이 포함되는데 제외해도 별 문제 없음

Sentinel: An Imaginary Leaf Node



(b) A red-black counterpart of (a)

(c) Implementation of (b)



✓ NIL leaf는 구현상 매우 유용하다

Sentinel is Useful

Programming example:

NIL is a **black**-colored **TreeNode** object

```
class TreeNode {  
    ... key;  
    TreeNode left;  
    TreeNode right;  
    int color;  
    ...  
}
```



if (t.left.color == black) ...

Theorem

* black height: 루트에서 리프 노드에 이르는 경로상에서
만나는 블랙 노드의 개수(루트는 제외)

Theorem

키가 총 n 개인 RB 트리의 가능한 최대 깊이는 $O(\log n)$ 이다

<Proof>

키의 총 수가 $n \rightarrow$ internal node의 수가 n .

\rightarrow 가장 이상적으로 균형잡힌 이진검색트리의 깊이는 $\lfloor \log_2 n \rfloor + 1$.

\rightarrow RB 트리가 이상적으로 만들어져도 black height는 $\lfloor \log_2 n \rfloor + 1$ 를 넘을 수 없다.

RB property ②에 의해,


루트에서 리프에 이르는 경로 상에서 레드 노드가 블랙 노드보다 많을 수 없다

RB 트리의 internal node의 경로 길이는 $2(\lfloor \log_2 n \rfloor + 1)$ 를 넘을 수 없다

이것은 $O(\log n)$ 이다.

Insertion

Do general BST insertion,
color **red** to the inserted node x ,
and link two NIL leaves from x




situation after insertion:
 p is **black** or **red**

Insertion

Situation after insertion: p is **black** or **red**

1. If p is **black**: satisfies all the RB properties. Completed!




There are only these four cases

Insertion

2. If p is **red**: RB property ② is violated

{ If p is the root, change p to **black**. Completed!
otherwise, repair (next page..)

root가 **black**이라는
제약을 없앤 뒤수습




There are only these two cases.
They are symmetric. Here I show just the left case.


Insertion

Two cases depending on p 's sibling s

Case 1: s is **red**




Case 2: s is **black**



Insertion

Case 1: s is **red**



{ If p^3 is **black**, completed
If p^3 is **red**, p^2 becomes new x : Recursive!



● : 색이 바뀐 노드

Change p and s to **black**, p^2 to **red**.

Insertion



Case 2: s is black

Case 2-1



✓ Case 2-2로

Case 2-2



✓ completed



Case 1: $O(\log n)$

Case 2: $\Theta(1)$

————→ $O(\log n)$ in total (considered only repairing)


Deletion

- We can **restrict to the cases** that the deleted node has
no child or **only one child**
 - reason: <쉽게 배우는 알고리즘>(p.174) or
<Introduction to Algorithms>(p.289)
 - m : 삭제될 노드
- If m is **red**: no problem! (m has no child)
- Even when m is black, no problem if m has a child(반드시 **red**다)!




Problematic Case

Problem occurs only when the black m has no child.




최초에는 왼쪽 박스와 같은 **NIL leaf** 밖에 가능하지 않으나
수선 과정에서 **recursive**하게 위로 올라가면서
이런 **tree**가 될 수도 있다

Classification of Cases




Case 1





Case 2

All Possible Cases


Case 1-1





Case 1-2




Case 1-3




Case 2-2



Case 2-3



Case 2-1




Case *-2

Case *-3


상관없다

Case 2-4




Five Groups


Case 1-1




Case 2-1




Case 2-4



Case 1-2




Case 2-2




Case *-2

Case 1-3



Case 2-3




Case *-3


Five Groups

⊖ : either **black** or **red**



Case 1-1




Case 2-1



Case 2-4



Case *-2





Case *-3

Repair Operations

⊖ : either **black** or **red**



● : node whose color changes or may change

Case 1-1





✓ completed

Case *-2



✓ completed



Case *-3



✓ Case *-2로



Repair Operations

Case 2-1



✓ p becomes new x : Recursive!

Case 2-4



✓ Case 1-1, 1-2, 1-3 중의 하나로

Case 2-1: $O(\log n)$

All the other cases except Case 2-1: $\Theta(1)$

————→ $O(\log n)$ in total (considered only repairing)

- ✓ In addition, intuitively think about why repairs (insertion and deletion) takes $O(\log n)$