



Hash Tables

Reminder:

Basics of Hash Tables

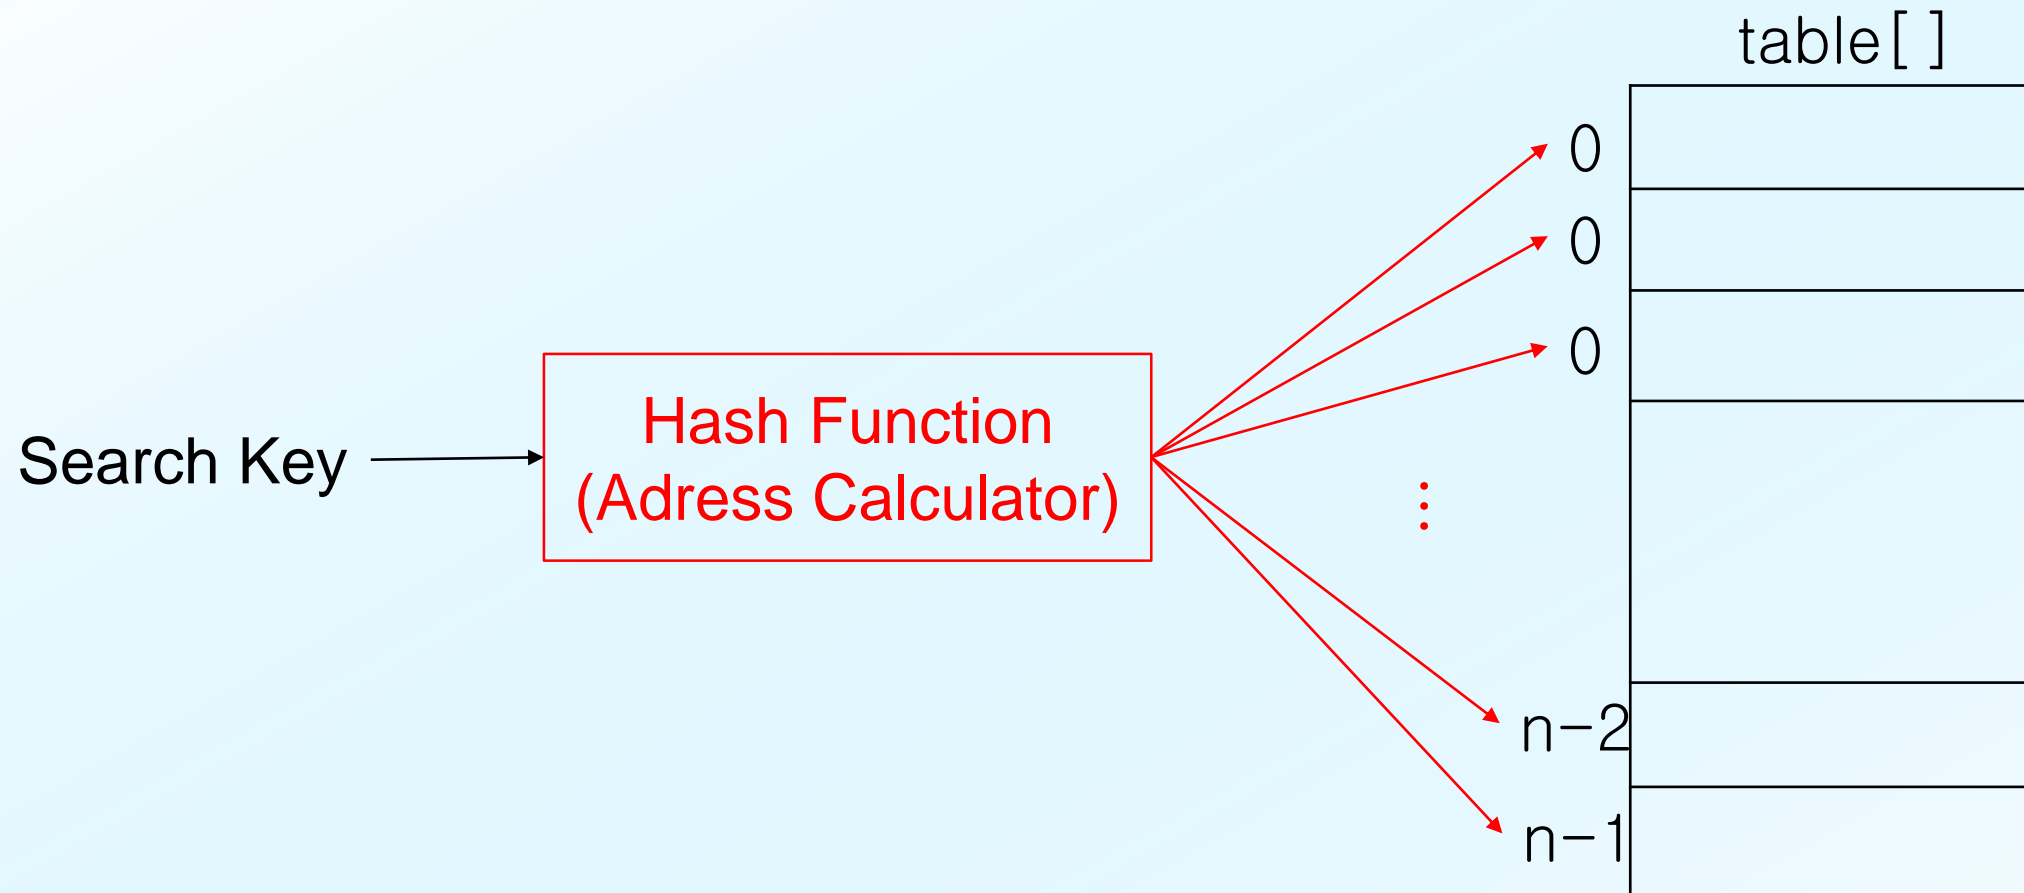
Want $\Theta(1)$ -Time Operations

- Array or linked list
 - Overall $O(n)$ time
- Binary search trees
 - Expected $\theta(\log n)$ -time search, insertion, and deletion
 - But, $\theta(n)$ in the worst case
- Balanced binary search trees
 - Guarantees $O(\log n)$ -time search, insertion, and deletion
 - Red-black tree, AVL tree
- Balanced k -ary trees
 - Guarantees $O(\log n)$ -time search, insertion, and deletion w/ smaller constant factor
 - 2-3 tree, 2-3-4 tree, B-trees
- Hash table
 - Expected $\theta(1)$ -time search, insertion, and deletion

Hash Tables

- Stack, queue, priority queue
 - do not support *search* operation
- Hash table support quick search, insertion, and deletion
 - But, does not support finding the minimum (or maximum) element
- Applications that need very fast operations
 - 119 emergent calls and locating caller's address
 - Air flight information system
 - 주민등록 시스템

Address Calculator



Hash Functions

- Toy functions
 - Selection digits
 - $h(001364825) = 35$
 - Folding
 - $h(001364825) = 1190$
- Modulo arithmetic
 - $h(x) = x \bmod \text{tableSize}$
 - *tableSize* is recommended to be prime
- Multiplication method
 - $h(x) = (xA \bmod 1) * \text{tableSize}$
 - *A*: constant in (0, 1)
 - *tableSize* is not critical, usually 2^p for an integer p


Collision Resolution

Collision:

a key maps to an occupied location in the hash table

$$h(224) = 224 \bmod 101 = 22$$

table[22] is
occupied



An example: $h(x) = x \bmod 101$

Collision resolution

- resolves collision by a seq. of hash values
- $h_0(x)(=h(x)), h_1(x), h_2(x), h_3(x), \dots$
- The most important in hash tables

Collision-Resolution Methods

Open addressing (resolves in the table)

- Linear probing

- $h_i(x) = (h_0(x) + i) \% \text{tableSize}$

- Quadratic probing

- $h_i(x) = (h_0(x) + i^2) \% \text{tableSize}$

- Double hashing

- $h_i(x) = (h_0(x) + i \cdot \beta(x)) \% \text{tableSize}$

- $\beta(x)$: another hash function

Full version:

$$h_i(x) = (h_0(x) + ai + b) \% \text{tableSize}$$

} Simple version

Full version:

$$h_i(x) = (h_0(x) + ai^2 + bi + c) \% \text{tableSize}$$

Separate chaining

- Each $\text{table}[i]$ is maintained by a linked list

Open Addressing


Linear probing

$$h_i(x) = (h_0(x) + i) \bmod \text{tableSize}$$

bad w/ primary clustering

Linear probing with

$$h_i(x) = (h_0(x) + i) \bmod 101$$



table[]		삽입 순서: 123, 24, 224, 22, 729, ...	
0			
	⋮		
22	123	$h_0(123)=h_0(224)=h_0(22)=h_0(729)=22$	
23	224	$i+1$	
24	24	$i+2$	$h_0(24) = 24$
25	22	$i+3$	
	729	$i+4$	
	⋮		
100			

Open Addressing

Quadratic probing

$$h_i(x) = (h_0(x) + i^2) \bmod \text{tableSize}$$

bad w/ secondary clustering

Quadratic probing with
 $h_i(x) = (h_0(x) + i^2) \bmod 101$



	table[]	
	⋮	
22	123	$i = 123 \bmod 101 = 22$
23	224	$i + 1^2$
	24	
26	22	$i + 2^2$
	⋮	
3	729	$i + 3^2$
1	⋮	

Open Addressing


Double hashing

$$h_i(x) = (h_0(x) + i\beta(x)) \bmod 101$$

Double hashing with

$$h_0(x) = x \bmod 101$$

$$\beta(x) = 1 + (x \bmod 97)$$



table[]		
	⋮	
22	123	$h_0(123) = h_0(224) = h_0(22) = h_0(729) = 22$
	⋮	
45	22	$\beta(22) = 23, h_1(22) = 45$
	⋮	
53	224	$\beta(224) = 31, h_1(224) = 53$
	⋮	
73	729	$\beta(729) = 51, h_1(729) = 73$
	⋮	

Be Careful in Deletion

Hash function:

$$h_i(x) = (h_0(x) + i) \bmod 13$$

0	13
1	1
2	15
3	16
4	28
5	31
6	38
7	7
8	20
9	
10	
11	
12	25

(a) Delete element 1

0	13
1	
2	15
3	16
4	28
5	31
6	38
7	7
8	20
9	
10	
11	
12	25

(b) Search 38,
wrong result!

0	13
1	DELETED
2	15
3	16
4	28
5	31
6	38
7	7
8	20
9	
10	
11	
12	25

(c) Okay: marking
with DELETED

Insertion

hashInsert(x):

◀ table[]: hash table, x : new key to insert

if (table[$h(x)$] is not occupied)

 table[$h(x)$] $\leftarrow x$

else

 Find an appropriate location k by a collision-resolution method

 table[k] $\leftarrow x$

numItems++

Deletion

hashDelete(x):

◀ table[]: hash table, x : key to delete

Find the location k of x by search

if (search was successful)

 table[k] \leftarrow DELETED

 numItems--

When the Load Factor is Higher than Wanted


$$\alpha = \frac{\text{\# of occupied slots}}{\text{hash table size}}$$

- A hash table performs bad when the load factor(α) is too high
- Generally, set a threshold and if the load factor surpasses it
 - Double the size of the hash table and
rehash all the elements in the table

Separate Chaining

Table[] is a header array of linked lists

No interference bet'n keys not collided
(Open addressing may interfere...)



Operations in Chained Hash Table

search(table[], x):
Search x in the list table[$h(x)$]

insert(table[], x):
Insert x in the list table[$h(x)$]

delete(table, x):
Delete x in the list table[$h(x)$]

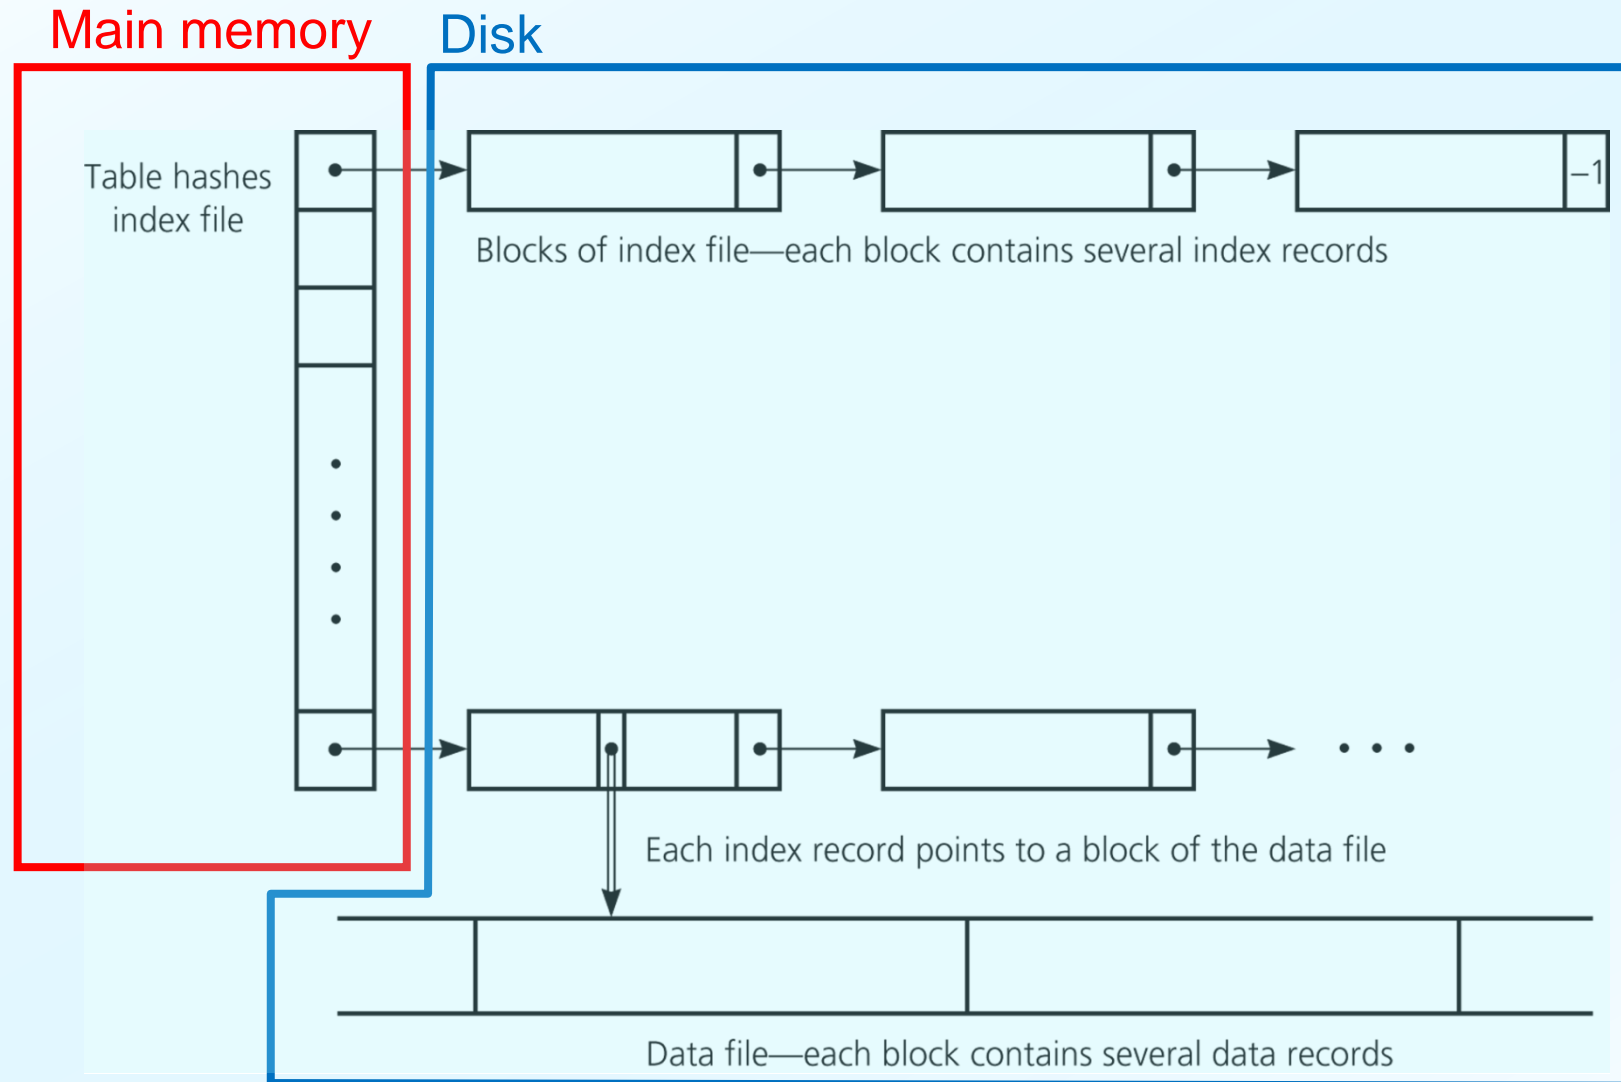
Observation

- No difference among probing methods when the load factor is low
- Successful search follows the same path as that of insertion

Internal/External Hashing

- Hash table is
 in the main memory(internal hashing) or
 in the disk(external hashing)
- In an external hashing, the # of disk accesses is critical

External Hash Table



Efficiency of Hash Tables

Search Time in Chaining

Assuming a uniform distribution of data,
a search takes $\Theta(\max(1, \alpha))$ on average

Search Time in Open Addressing

Assumption (**uniform hashing**)

- $h_0(x), h_1(x), \dots, h_{m-1}(x)$ is a permutation of $\{0, 1, \dots, m-1\}$
- Every permutation is equally likely

Note: collision probability = $\frac{n}{m}$
{ m: table size
n: # of elements in the table (= # of occupied slots)

[Theorem 1]

The expected #probes in an unsuccessful search or an insertion is at most $\frac{1}{1-\alpha}$

<proof>

$p_i = \text{Pr}(\text{exactly } i \text{ probes access occupied slots})$

$q_i = \text{Pr}(\text{at least } i \text{ probes access occupied slots})$

$$\begin{aligned} \text{Expected \# probes} &= 1 + \sum_{i \geq 1} i p_i \\ &= 1 + \sum_{i \geq 1} i (q_i - q_{i+1}) \\ &= 1 + \sum_{i \geq 1} q_i \\ &\leq 1 + \sum_{i \geq 1} \alpha^i \longleftarrow q_i = \frac{n}{m} \frac{n-1}{m-1} \cdots \frac{n-i+1}{m-i+1} \leq \left(\frac{n}{m}\right)^i = \alpha^i \\ &= \frac{1}{1-\alpha} \end{aligned}$$

[Theorem 2]

The expected #probes in a successful search
is at most $\frac{1}{\alpha} \ln \frac{1}{1-\alpha}$

Note: a successful search
exactly follows the path of insertion

<proof>

- The load factor α right after i^{th} key had been inserted was $\frac{i}{m}$
- If x is the $(i + 1)^{th}$ key inserted, then the expected #probes in a successful search for x is, by the previous thm, at most $\frac{1}{1-\frac{i}{m}}$
- Average over all keys

$$\begin{aligned} \frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} &= \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} \\ &\leq \frac{1}{\alpha} \int_0^n \frac{1}{m-x} dx \\ &= \frac{1}{\alpha} \ln \frac{1}{1-\alpha} \end{aligned}$$

A Creative Utilization of Hash Tables

Minhash

- Suggested by Andrei Broder, 1997
- Min-wise locality sensitive permutation hashing
- Fast computation of similarity of two sets is possible


Similarity of two {
vectors
documents
web pages
stock patterns
...

Jaccard Similarity

Jaccard
similarity


Two sets A, B

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$



e.g. $A = \{1, 2, 3, 4, 5\}$
 $B = \{3, 4, 5, 6, 7\}$ $\rightarrow J(A, B) = \frac{3}{7}$

For sets A_1, A_2, \dots, A_n ,
we often need to compute their pairwise similarities
or similarity to another set B



$h_{\min}(S)$: A Permutation Hashing

$h(x)$: a hash function

$h_{\min}(S) = x \in S$ that minimizes $h(x)$


$S = \{a, b, c, d, e\}$

$h(a) h(b) h(c) h(d) h(e)$

minimum

Then, $h_{\min}(S) = d$

example



$$\text{Prob}(h_{\min}(A) = h_{\min}(B)) = J(A, B)$$

Usage in Fields

Using one $h_{min}()$ just probabilistically matches with Jaccard similarity

Prepare many enough $h_{min}()$'s: $h_{min}^1(), h_{min}^2(), \dots, h_{min}^k()$ ← k different hash functions


For all $A_i, i = 1, 2, \dots, n$, compute (just one time) $h_{min}^1(), h_{min}^2(), \dots, h_{min}^k()$

→ $J(A_i, A_j) = \frac{\text{\# of the same } h_{min}'s}{k}$


$= \frac{\sum_{r=1}^k \delta(h_{min}^r(A_i), h_{min}^r(A_j))}{k},$

$$\delta(a, b) = \begin{cases} 1, & \text{if } a = b \\ 0, & \text{if } a \neq b \end{cases}$$

An example of Monte Carlo approximation
(random sampling based...)



Applying Minhash to DNA Pairwise Similarity



$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} \approx \frac{|S(A \cup B) \cap S(A) \cap S(B)|}{|S(A \cup B)|}$$

Overview of the MinHash bottom sketch strategy for estimating the Jaccard index. First, the sequences of two datasets are decomposed into their constituent k-mers (*top, blue and red*) and each k-mer is passed through a hash function h to obtain a 32- or 64-bit hash, depending on the input k-mer size. The resulting hash sets, A and B , contain $|A|$ and $|B|$ distinct hashes each (*small circles*). The Jaccard index is simply the fraction of shared hashes (*purple*) out of all distinct hashes in A and B . This can be approximated by considering a much smaller random sample from the union of A and B . MinHash sketches $S(A)$ and $S(B)$ of size $s = 5$ are shown for A and B , comprising the five smallest hash values for each (*filled circles*). Merging $S(A)$ and $S(B)$ to recover the five smallest hash values overall for $A \cup B$ (*crossed circles*) yields $S(A \cup B)$. Because $S(A \cup B)$ is a random sample of $A \cup B$, the fraction of elements in $S(A \cup B)$ that are shared by both $S(A)$ and $S(B)$ is an unbiased estimate of $J(A, B)$.

GTCGGATACCCAGCCCGTGTGAGGCTCCCTCGACGAGTCGAGTAA
GTCGGAAACCCAGCCCATGTGGAAGCTCCCTCGACGAGTCGAGTAA
GTAGGAAACCATGTCTCATGTGGAAGCTCCCTCGACGAGTCGAGTAA
GTAGGAAACCATGTCTCATGTGGAAGCTCCCTCGACGAGTCGAGTAA
GTAGGAAACCATGTCTCATGTGGAAGCTCCCTCGACGAGTCGAGTAA
GTAGGAAACCATGTCTCATGTGGAAGCTCCCTCGACGAGTCGAGTAA
GTAGGACACCCAGCCCGTGTGACGCTCCCTCGACGAGTCGAGTAA
GTAGGATACCCAGCCCATGTGACGCTCCCTCGACGAGTCGAGTAA
GTAGGTTACCCAGCCCGTGTGAGGCTCCCTCGACGAGTCGAGTAA
GTAGGTTACGACGCCCGTGTGAGGCTCCCTCGACGAGTCGAGTAA
GTAGGATACCCAGCCCATGTGGAAGCTCCCTCGACGAGTCGAGTAA

A good example,
although they made some variation